

Virtual Machine Security

CSE443 - Spring 2012

Introduction to Computer and Network Security

Professor Jaeger

www.cse.psu.edu/~tjaeger/cse443-s12/


Operating System Quandary

- Q: What is the primary goal of system security?
 - OS enables multiple users/programs to share resources on a physical device
- Q: What happens when we try to enforce Mandatory Access Control policies on UNIX systems
 - Think SELinux policies
- What can we do to simplify?



Virtual Machines

- Instead of using system software to enable sharing, use system software to enable **isolation**
- Virtualization
 - “a technique for hiding the physical characteristics of computing resources from the way in which others systems, applications, and end users interact with those resources”
- Virtual Machines
 - Single physical resource can appear as multiple logical resources



Virtual Machine Architectures


- ***Full system simulation***
 - CPU can be simulated
- ***Paravirtualization (Xen)***
 - VM has a special API
 - Requires OS changes
- ***Native virtualization (VMWare)***
 - Simulate enough HW to run OS
 - OS is for same CPU
- ***Application virtualization (JVM)***
 - Application API



- ***Type I***
 - Lowest layer of software is VMM
 - E.g., Xen, VAX VMM, etc.
- ***Type II***
 - Runs on a host operating system
 - E.g., VMWare, JVM, etc.
- Q: What are the trust model issues with Type II compared to Type I?


Virtual Machine Types

Type 2 VMM




JVM
CLR
VMware Workstation

Hybrid VMM



MS Virtual Server
KVM


Type 1 VMM



VMware ESX
Xen
MS Hyper-V

VM Security

- Isolation of VM computing
- Like a separate machine



- First system design to examine virtualization in the context of information flow security
- Virtualization mechanisms necessary to implement a *reference validation mechanism* that satisfies the *reference monitor concept*
- Assure system design and implementation to the highest level
 - *A1 level per the Orange Book*
- Control all system information flows according to *MLS* and *Biba integrity* policies (modulo exceptions in “privileges”)
- Also, *covert channel countermeasures* were produced, approximating noninterference
- System was piloted, but not released commercially

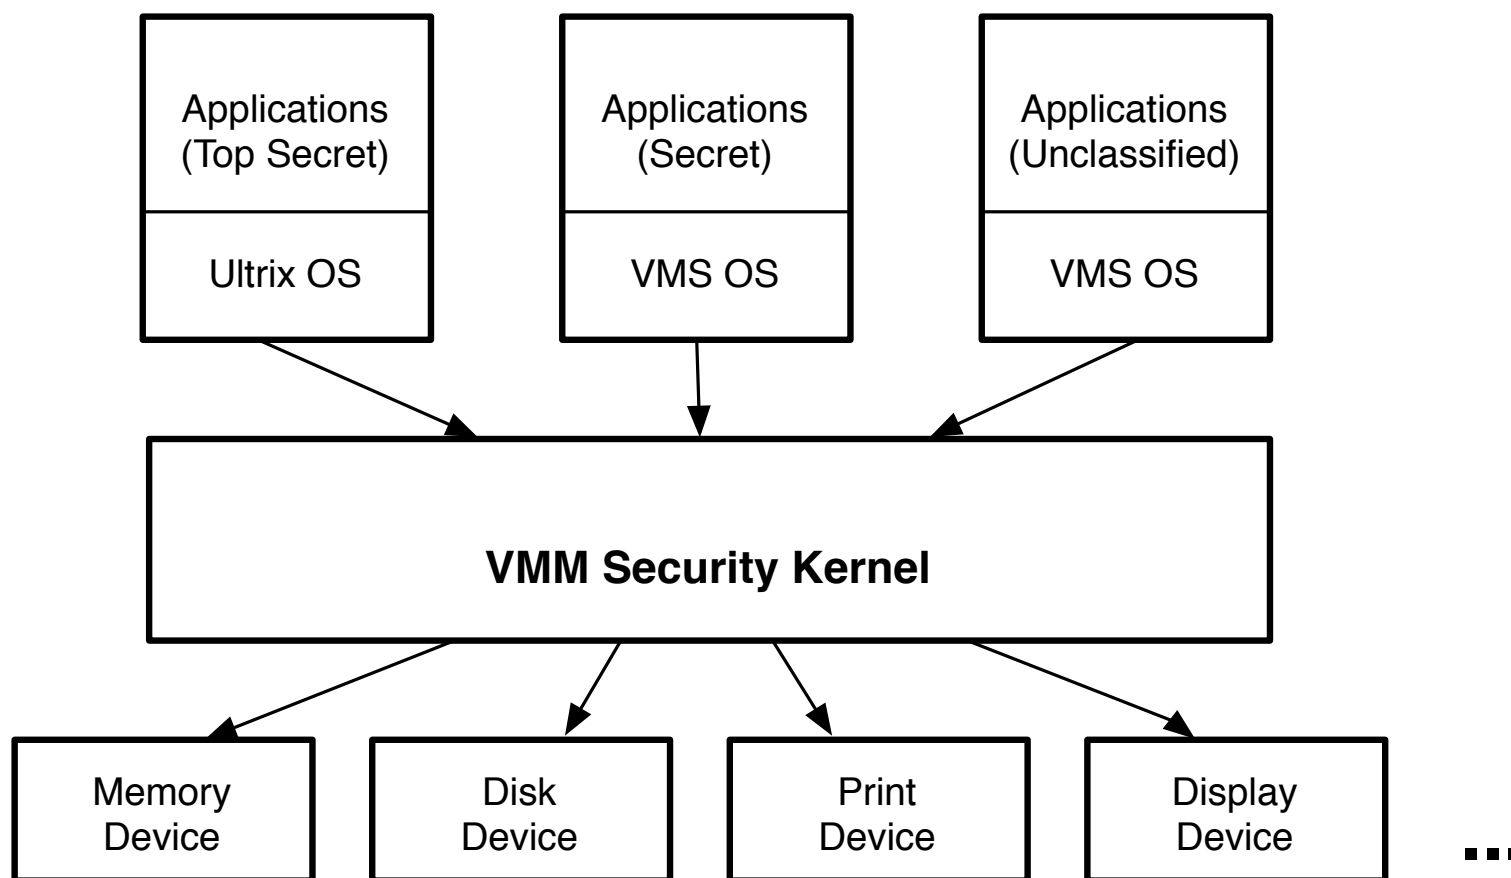
- Key design tasks of secure VMM
 - ▶ Virtualize processor
 - All security-sensitive instructions must be mediated by VMM
 - ▶ VMM protection ring
 - VMM must be deployed in a more privileged protection ring than the VMs
 - ▶ I/O emulation
 - Privileged I/O tasks must be executed in VMM or trusted VM
 - ▶ Self-virtualizable
 - OS must not detect when running on a VMM (or VMMs)

- Security-Sensitive Instructions
 - ▶ *Instructions that read or modify privileged system state*
- Privileged Instructions
 - ▶ *Instructions that cause a trap when executed in a non-privileged ring*
- All security-sensitive instructions must be privileged to enable the VMM to manage privileged system state (rather than individual VMs)
- This requirement was not met by VAX hardware nor x86 originally

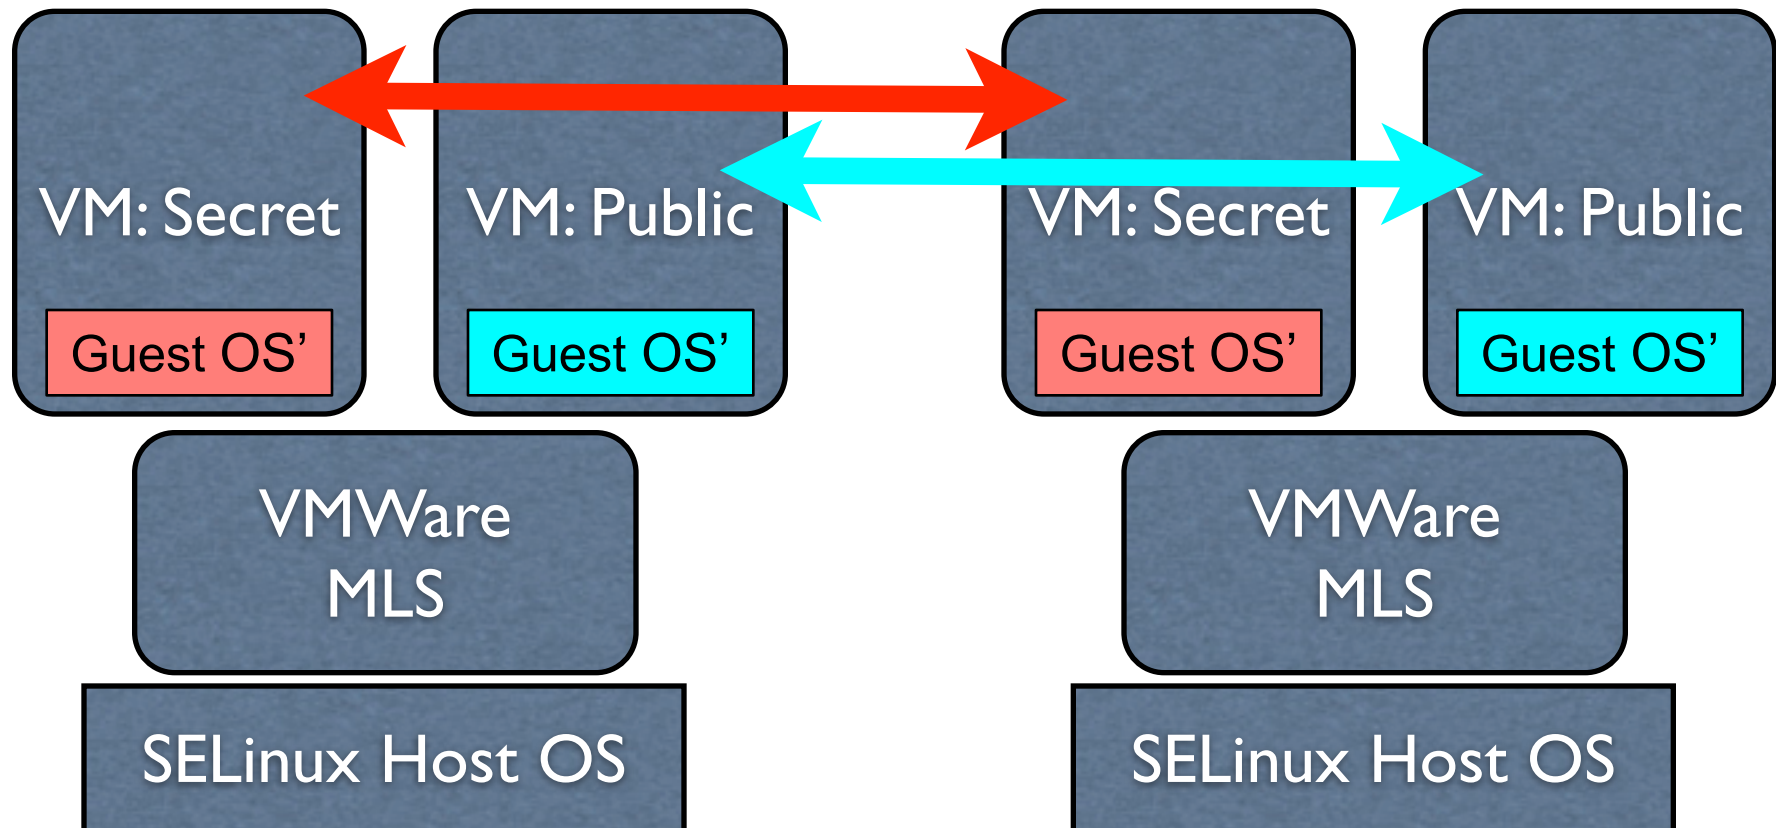
- Access to devices is expected by each operating system, but this access is security-sensitive
 - Thus, devices are virtualized
- Access to devices must be directed to the party with physical device access
 - Memory-mapped I/O uses unprivileged instructions
- **VAX VMM adds a layer of indirection**
 - I/O interface that causes a trap
 - OS must be modified to use that interface (paravirtualize)

- **Driver management**
 - ▶ In VAX VMM, all drivers were in the VMM kernel
 - ▶ This was for assurance, but added code to VMM
 - Drivers are outside the VMM in most systems
- **DMA**
 - ▶ Devices can use this mechanism to write to physical memory, but under guidance of untrusted VMs
 - VAX VMM trusted drivers, but not practical today
- **Performance** – E.g., page table lookups

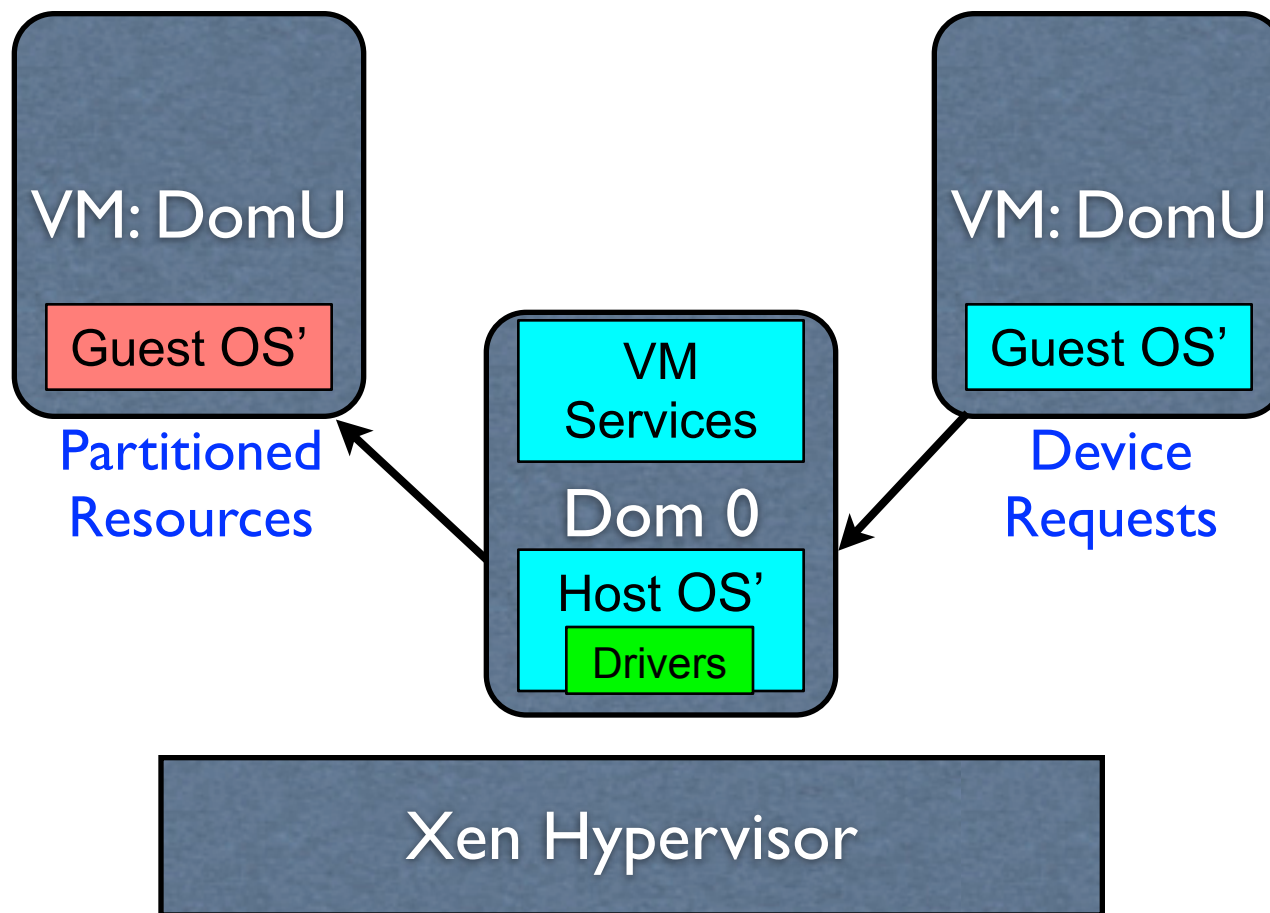
VAX VMM System



- Isolated networks of VMs
- Alternative to “air gap” security

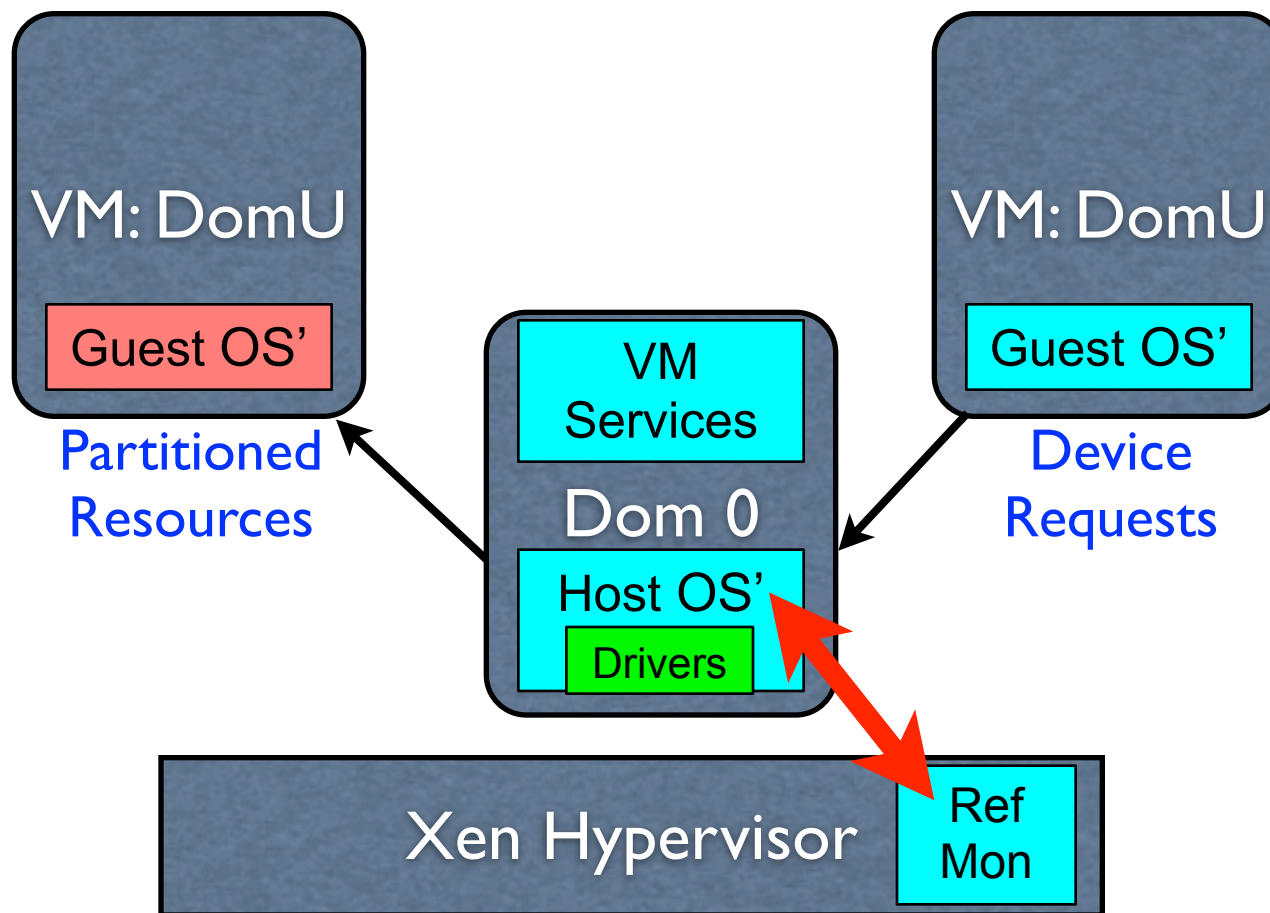


- Privileged VM



Xen sHype


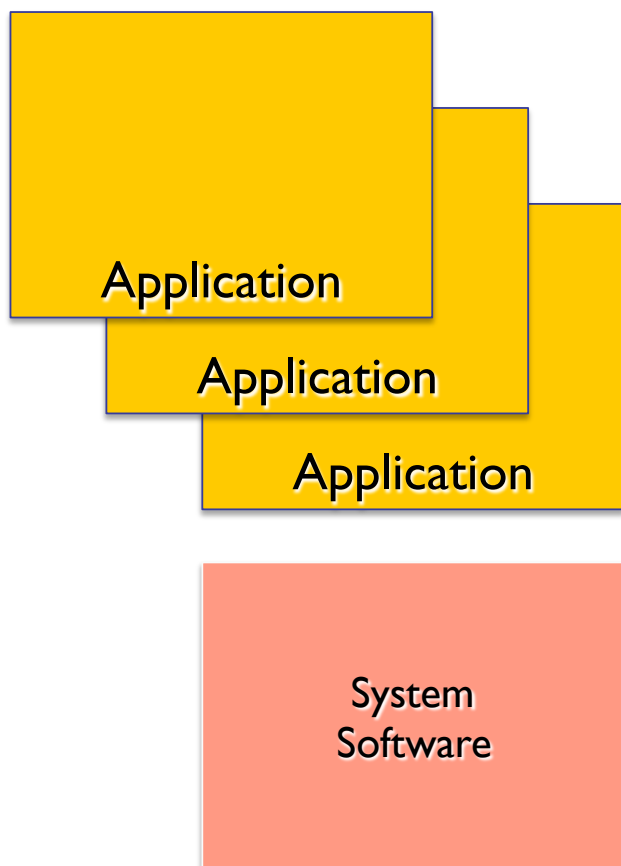
- Controlled information flows among VMs



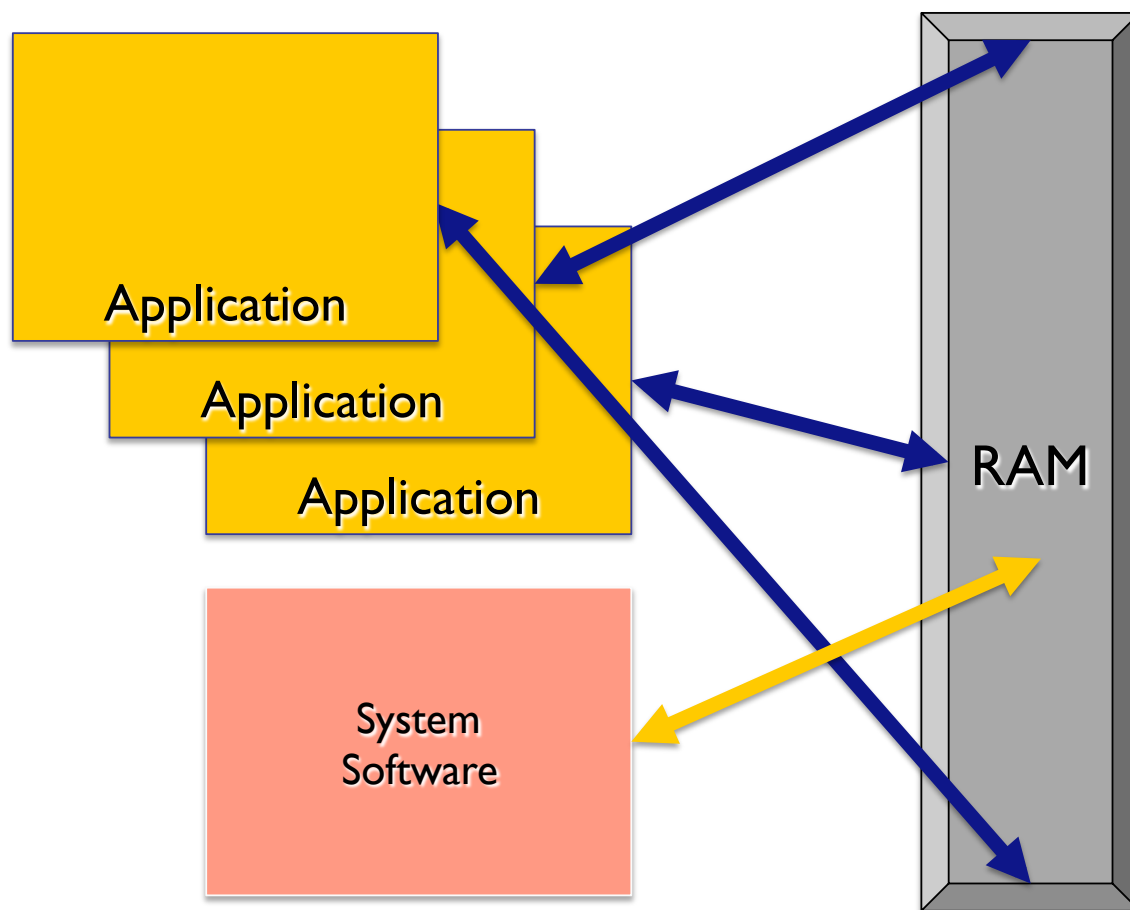
- Type Enforcement over VM communications
 - VM labels are subjects
 - VM labels are objects
- How do VMs communicate in Xen?
 - **Grant tables**: pass pages between VMs
 - **Event channels**: notifications (e.g., when to pass pages)
- sHype controls these
- Q: What about VM communication across systems?

- Comprehensive Reference Monitor interface for Xen
 - Based on LSM ideas
- Includes about 57 “hooks” (more expected)
 - Supports sHype hooks
 - Plus, hooks for VM management, resource partitioning
- Another aim: Decompose domain 0
 - Specialize kernel for privileged operations
 - E.g., Remove drivers

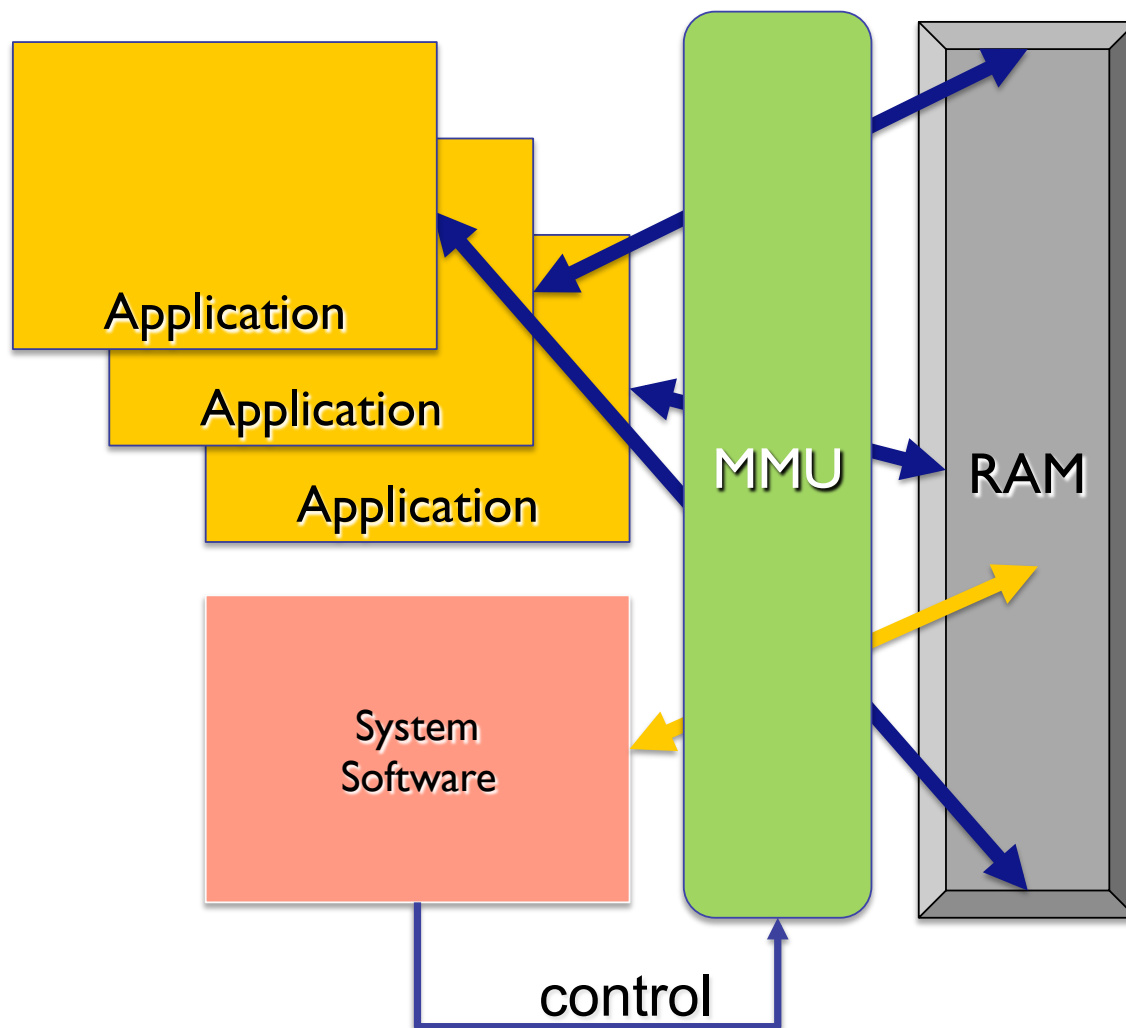
IOMMU Role in the System



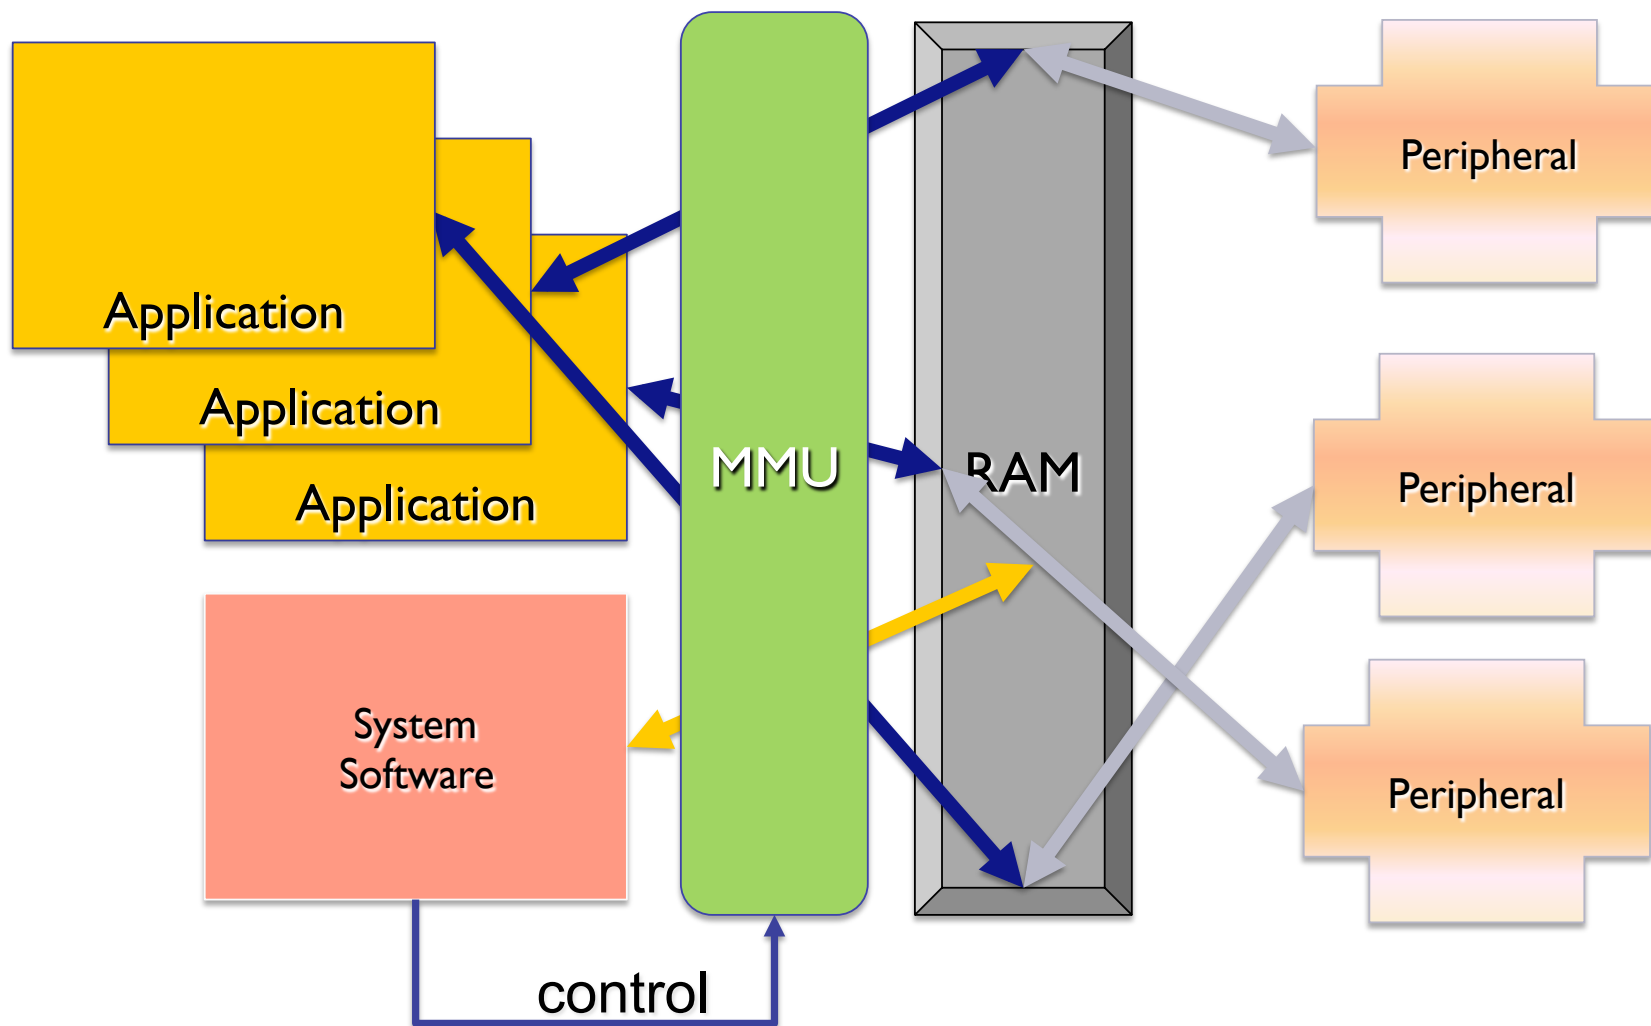
IOMMU Role in the System



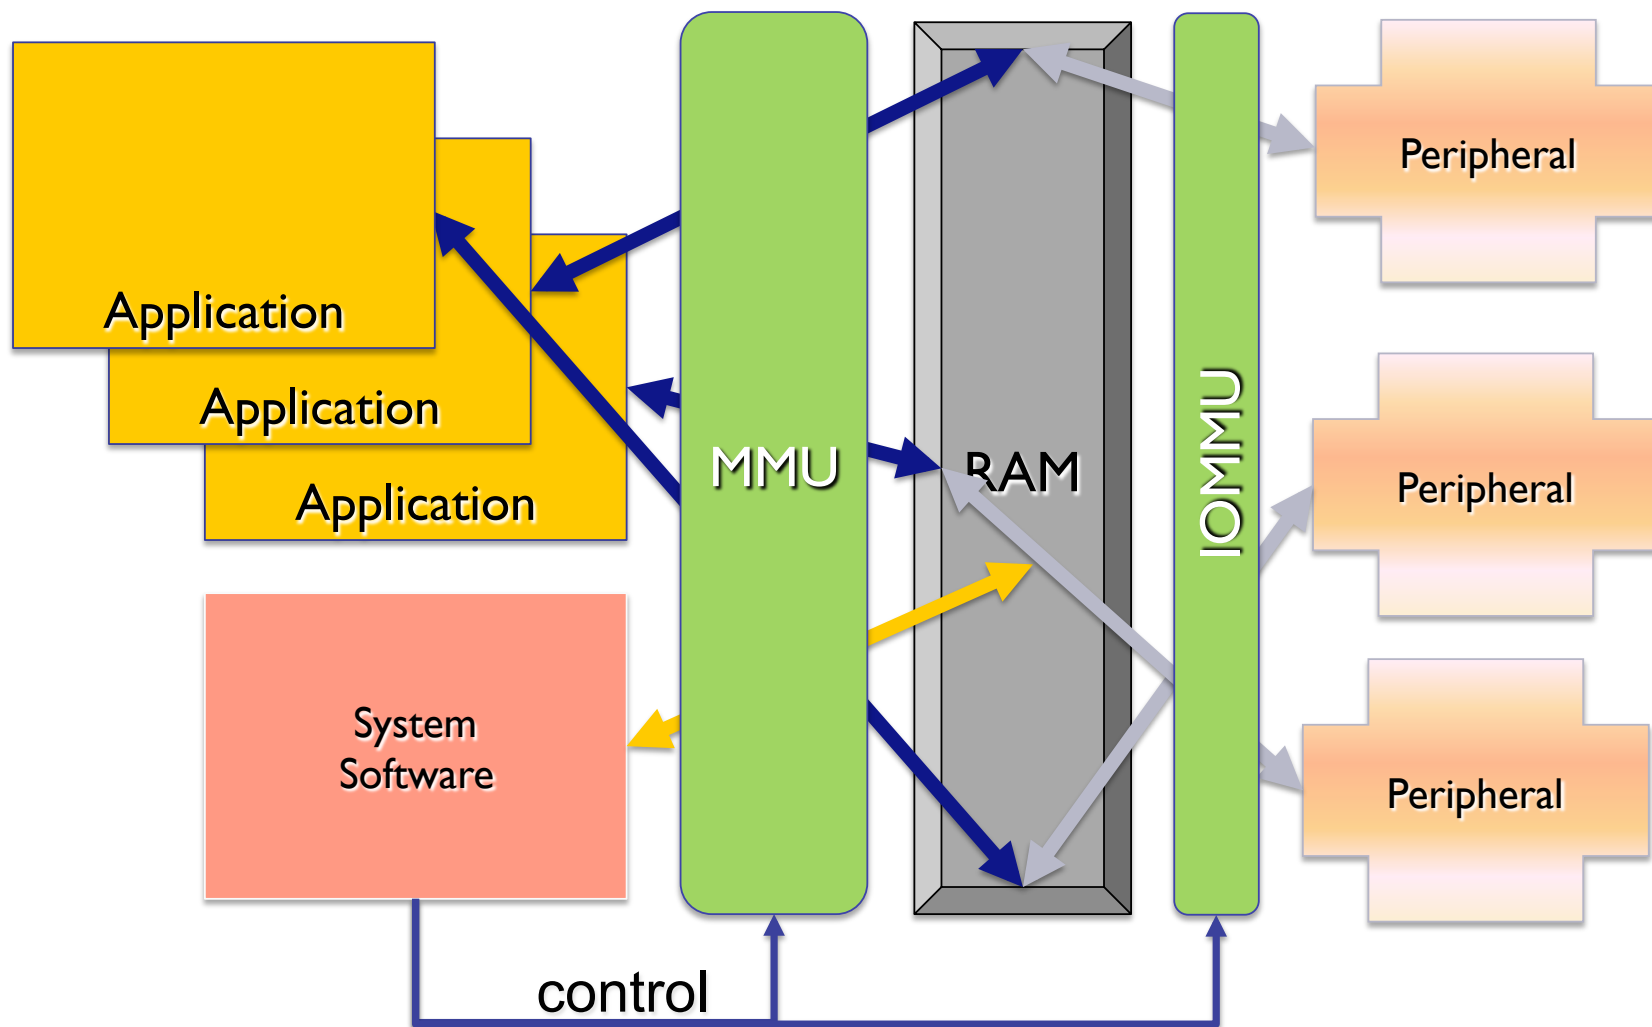
IOMMU Role in the System



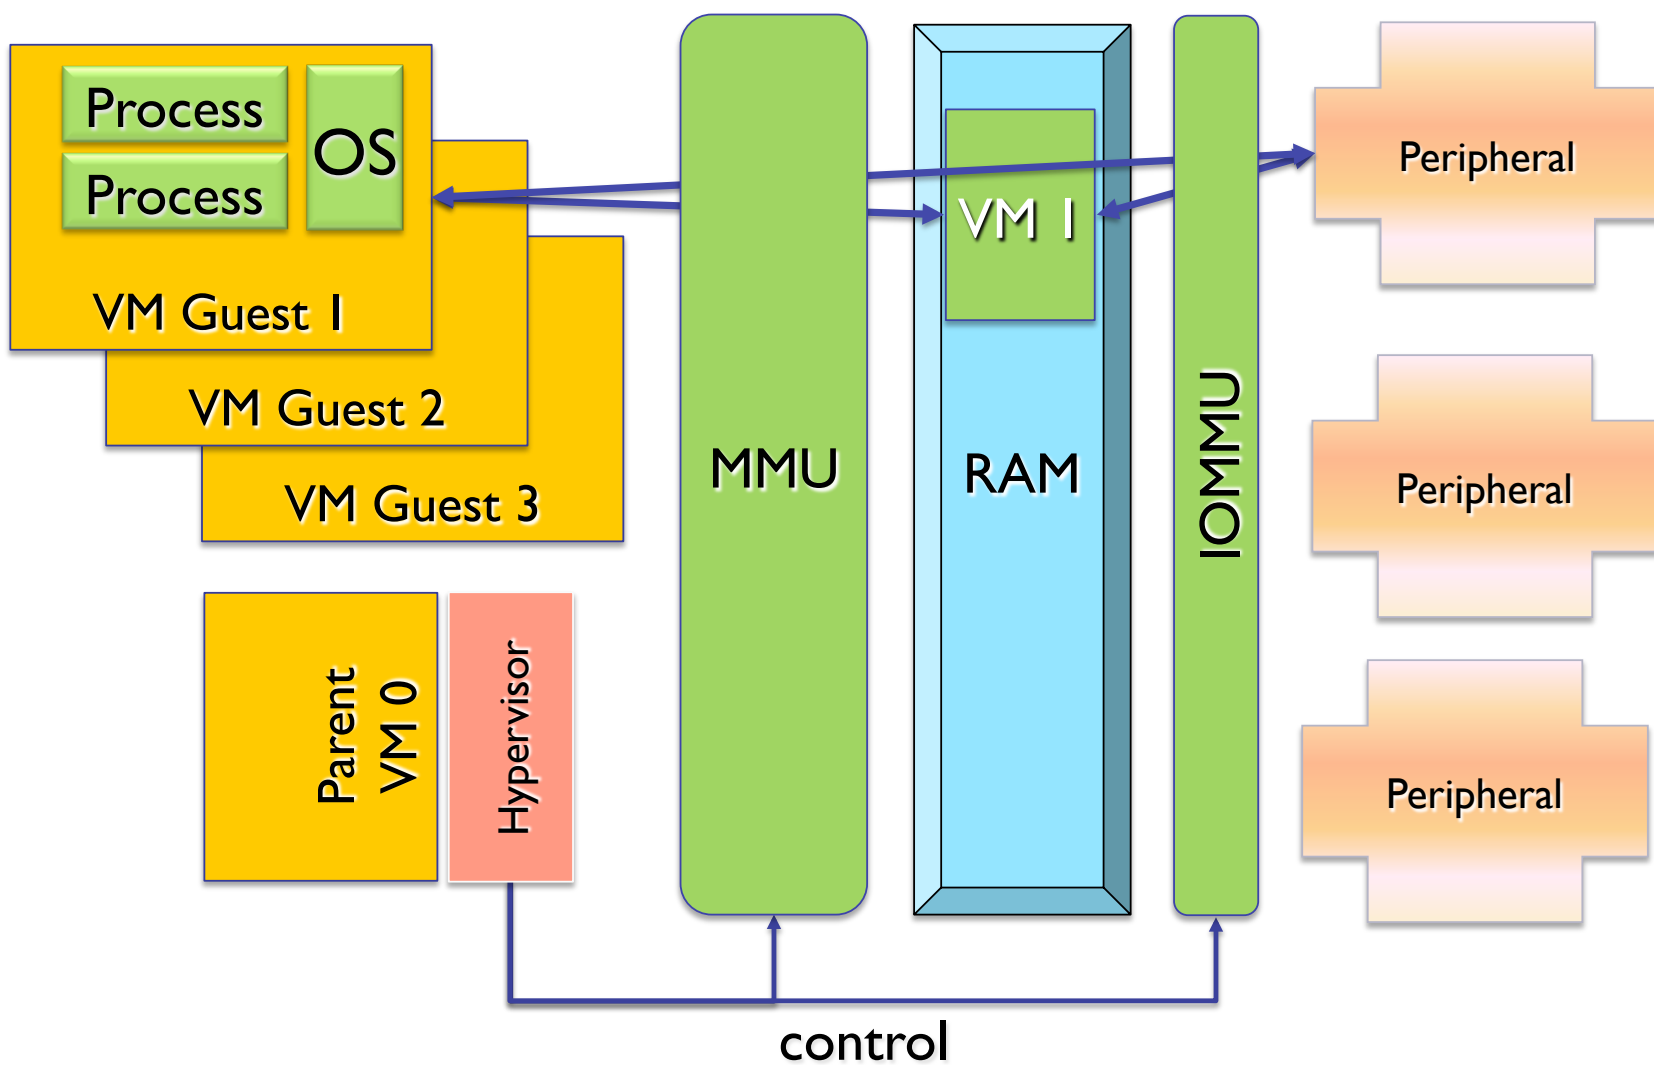
IOMMU Role in the System



IOMMU Role in the System



I/O Device Assignment



- Aim is simplicity
 - Are we achieving this?
- Do we care what happens in the VMs?
 - When might we care?
- Trusted computing base
 - How does this compare to traditional OS?

- How does the insertion of a virtual machine layer change the threats against the system?

- Rootkit
 - Malicious software installed by an attacker on a system
 - Enable it to run on each boot
- OS Rootkits
 - Kernel module, signal handler, ...
 - When the kernel is booted, the module is installed and intercepts user process requests, interrupts, etc.
 - E.g., keylogger
- VM Rootkit
 - Research project from Michigan and Microsoft
 - If security service runs in VM, then a rootkit in VMM can evade security
 - E.g., Can continue to run even if the system appears to be off

Take Away

- VM systems focus on isolation
 - Enable reuse, but limited by security requirements
- Enable limited communication
 - The policies are not trivial, but refer to coarser-grained objects

